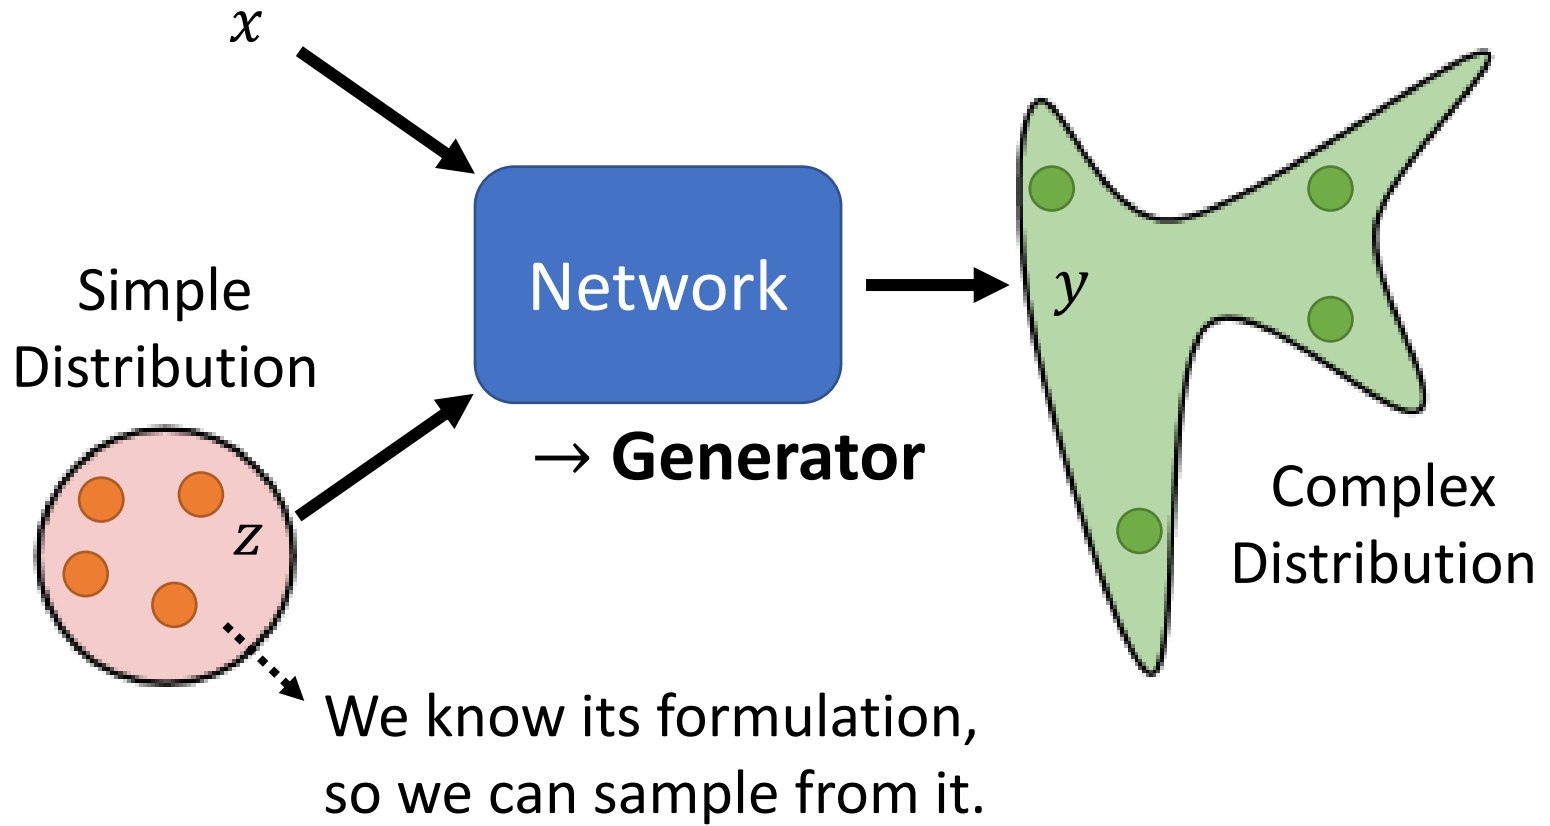





Generation

Hung-yi Lee 李宏毅

Network as Generator




Why distribution?



Real Video


Video Prediction




Previous
frames



next
frame




Why distribution?




Prediction

Video Prediction



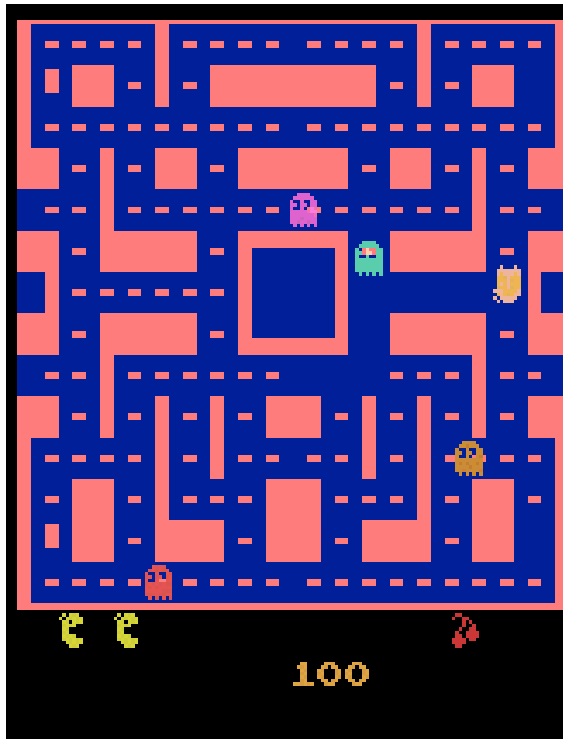
Previous
frames



turn
right


turn
left

Why distribution?




Prediction

Video Prediction



Previous frames



Why distribution?

(The same input has different outputs.)

- Especially for the tasks needs “*creativity*”

Drawing

Character
with red eyes



Chatbot

你知道輝夜是
誰嗎？



她是秀知院學生會 ...

她開創了忍者時代 ...



Generative Adversarial Network (GAN)

GAN

- How to pronounce “GAN”?



Google 小姐

All Kinds of GAN ...

<https://github.com/hindupuravinash/the-gan-zoo>

GAN

ACGAN

BGAN

CGAN

DCGAN

EBGAN

fGAN

GoGAN

⋮

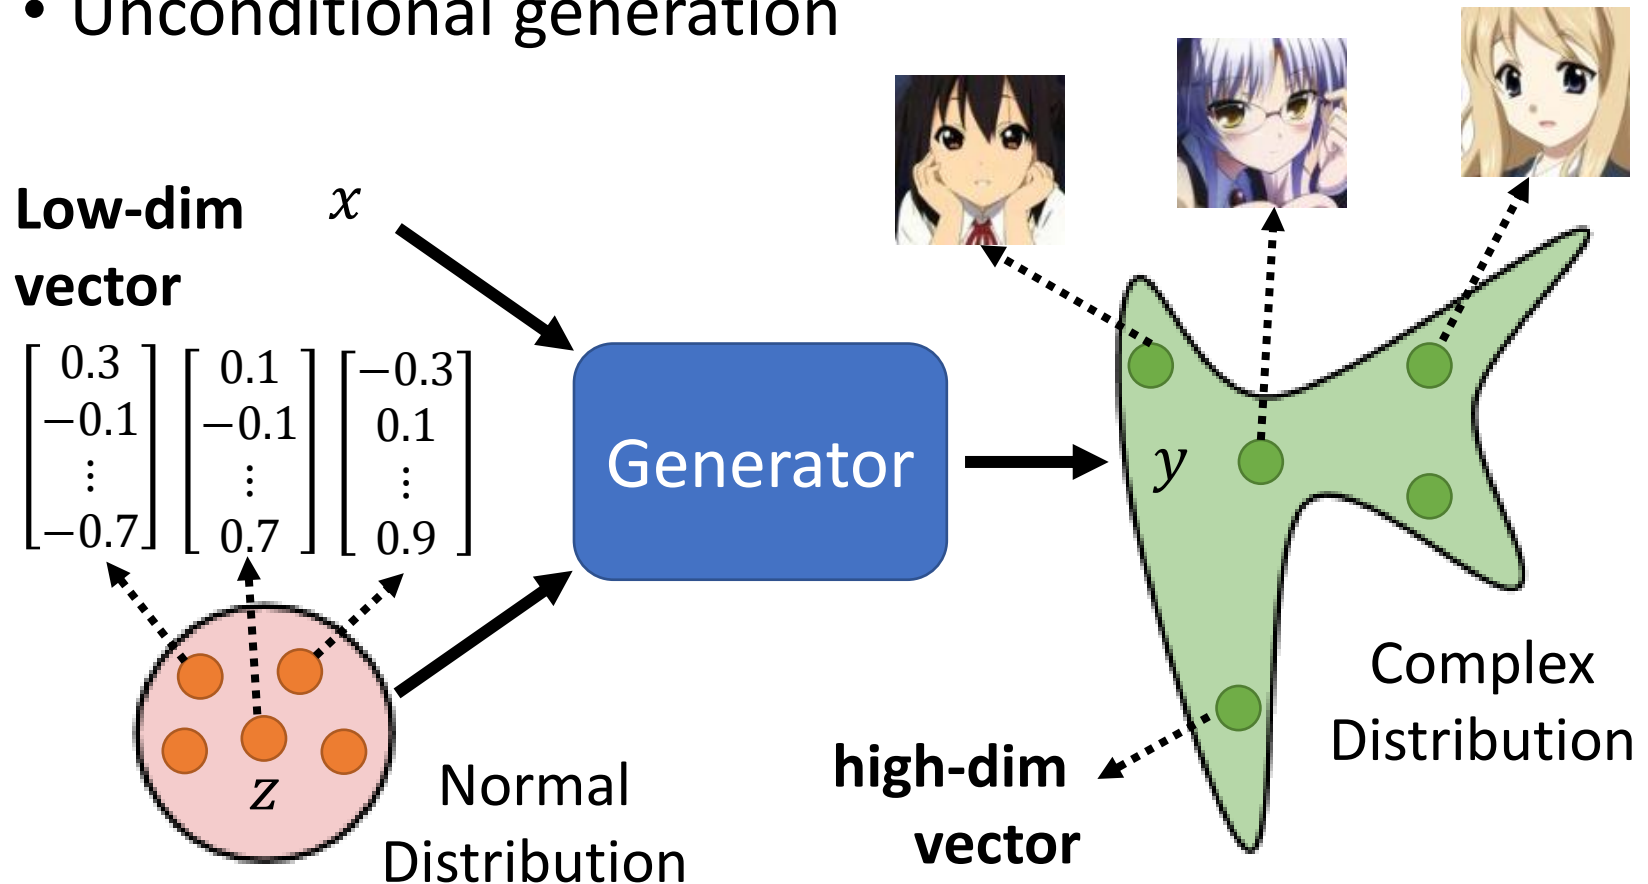
- SeUDA - [Semantic-Aware Generative Adversarial Nets for Unsupervised Domain Adaptation and Image Segmentation](#)
- SG-GAN - [Semantic-aware Grad-GAN for Virtual-to-Real Urban Scene Adaption](#) (github)
- SG-GAN - [Sparsely Grouped Multi-task Generative Adversarial Networks for Facial Attribute Transfer](#)
- SGAN - [Texture Synthesis with Spatial Generative Adversarial Networks](#)
- SGAN - [Stacked Generative Adversarial Networks](#) (github)
- SGAN - [Steganographic Generative Adversarial Networks](#)
- SGAN - [SGAN: An Alternative Training of Generative Adversarial Networks](#)
- SGAN - [CT Image Enhancement Using Stacked Generative Adversarial Networks and Total Variation](#) (github)
- sGAN - [Generative Adversarial Training for MRA Image Synthesis Using Multi-Contrast Self-Supervised Loss](#)
- SiftingGAN - [SiftingGAN: Generating and Sifting Labeled Samples to Improve the Remote Sensing Image Classification Baseline in vitro](#)
- SiGAN - [SiGAN: Siamese Generative Adversarial Network for Identity-Preserving Face Image Generation](#)
- SimGAN - [Learning from Simulated and Unsupervised Images through Adversarial Training](#)
- SisGAN - [Semantic Image Synthesis via Adversarial Learning](#)

Mihaela Rosca, Balaji Lakshminarayanan, David Warde-Farley, Shakir Mohamed, "Variational Approaches for Auto-Encoding Generative Adversarial Networks", arXiv, 2017

²We use the Greek α prefix for α -GAN, as AEGAN and most other Latin prefixes seem to have been taken
<https://deephunt.in/the-gan-zoo-79597dc8c347>.


Anime Face Generation

- Unconditional generation




Discriminator

It is a neural network
(that is, a function).

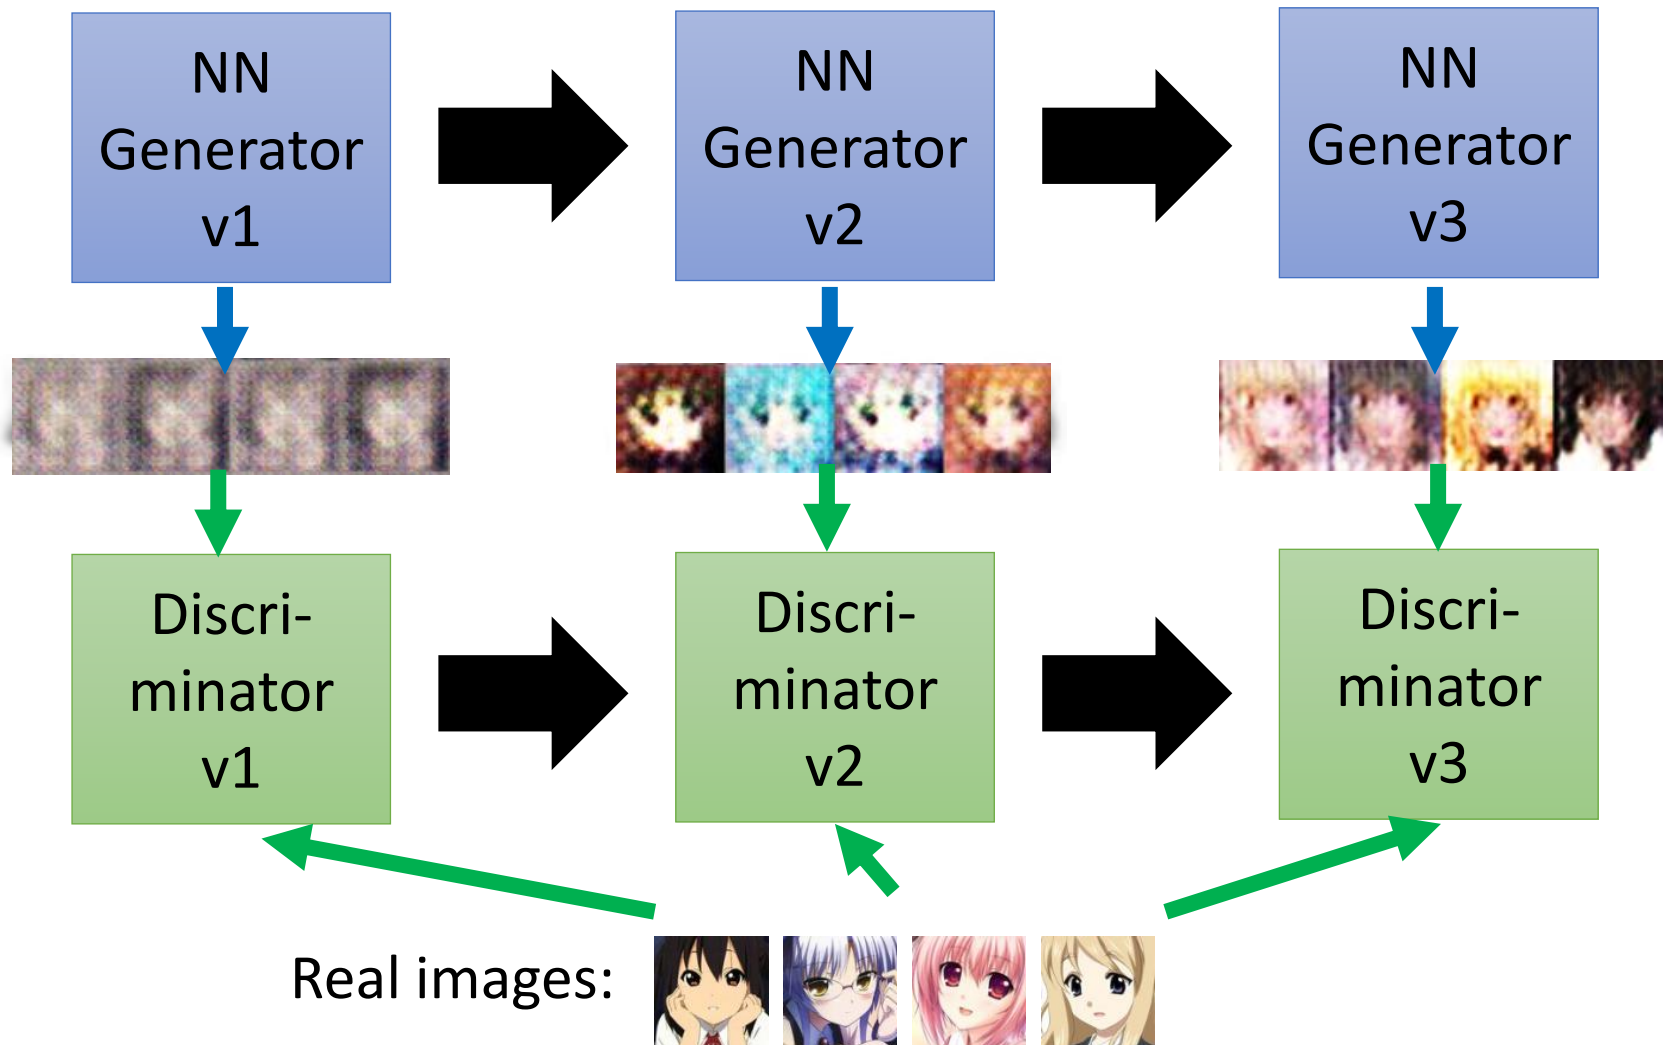


Basic Idea of GAN



Basic Idea of GAN

This is where the term
“*adversarial*” comes from.

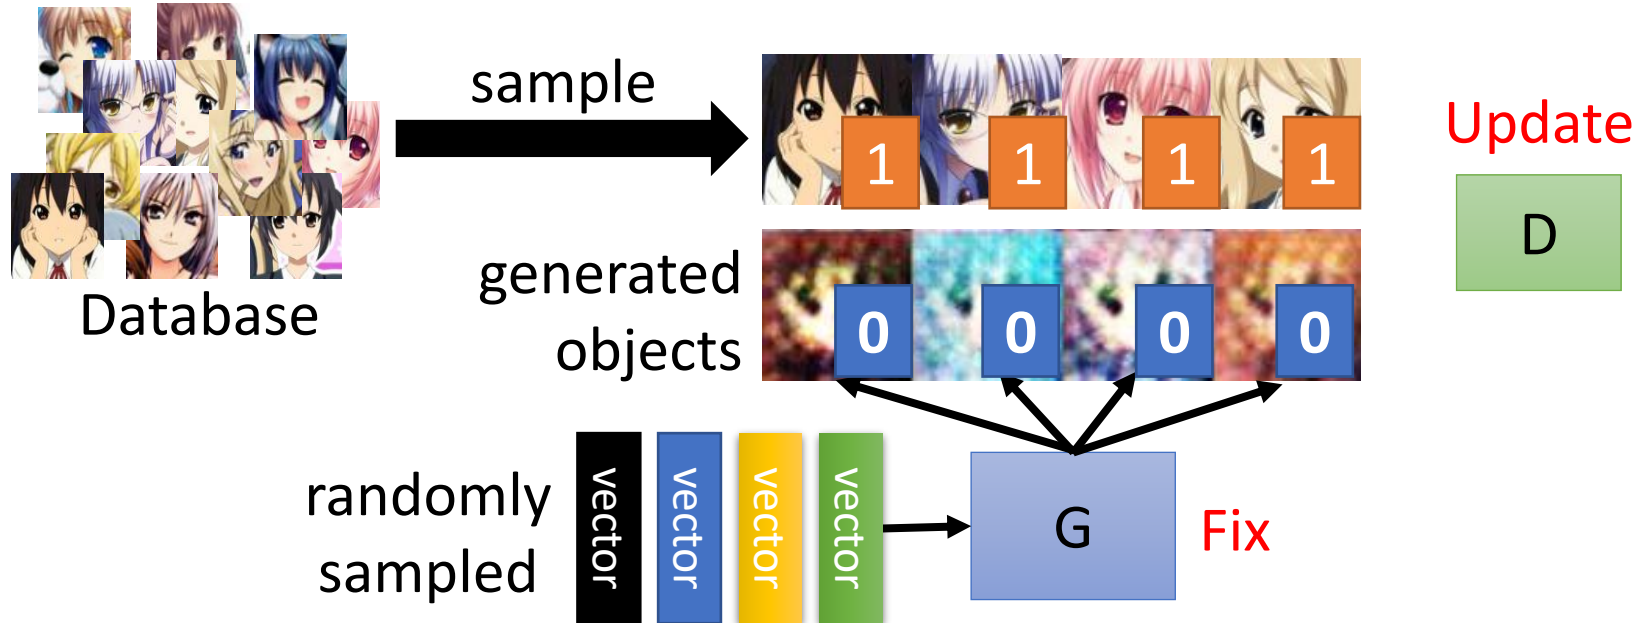


Algorithm

- Initialize generator and discriminator
- In each training iteration:



Step 1: Fix generator G, and update discriminator D



Discriminator learns to assign high scores to real objects and low scores to generated objects.


Algorithm

- Initialize generator and discriminator
- In each training iteration:



Step 2: Fix discriminator D, and update generator G

Generator learns to “fool” the discriminator




Algorithm


- Initialize generator and discriminator
- In each training iteration:



Sample some
real objects:



Generate some
fake objects:




fix

Update




update



fix



Anime Face Generation



100 updates

Source of training data: <https://zhuanlan.zhihu.com/p/24767059>

Anime Face Generation



1000 updates

Anime Face Generation



2000 updates

Anime Face Generation



5000 updates

Anime Face Generation



10,000 updates

Anime Face Generation



20,000 updates

Anime Face Generation



50,000 updates



The faces
generated by
machine.

圖片生成：
吳宗翰、謝濬丞、
陳延昊、錢柏均

In 2019, with StyleGAN




Source of video:

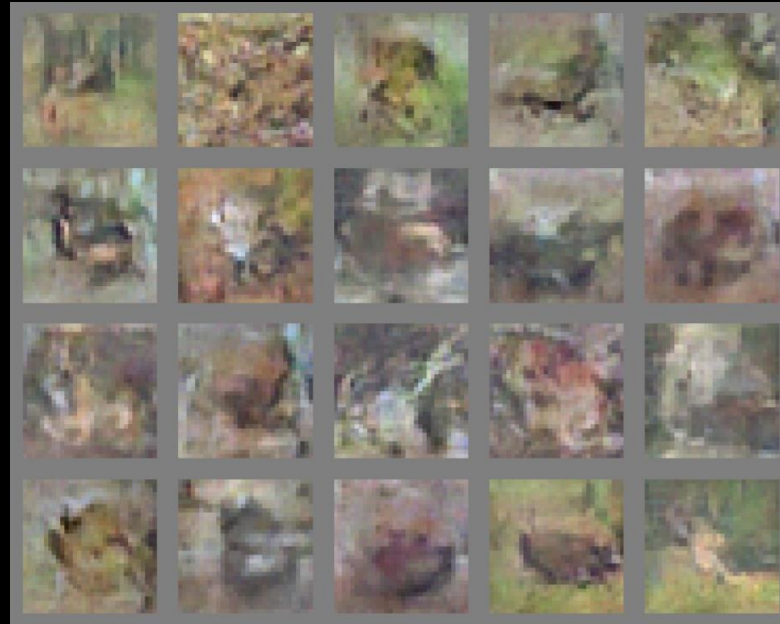
<https://www.gwern.net/Faces>



Progressive GAN |

<https://arxiv.org/abs/1710.10196>





The first GAN

<https://arxiv.org/abs/1406.2661> (Ian J. Goodfellow)



Today BigGAN

<https://arxiv.org/abs/1809.11096>




Theory behind GAN

$$\text{c.f. } w^*, b^* = \arg \min_{w, b} L$$

Our Objective

Normal
Distribution



$$G^* = \arg \min_G \underline{Div}(P_G, P_{data})$$

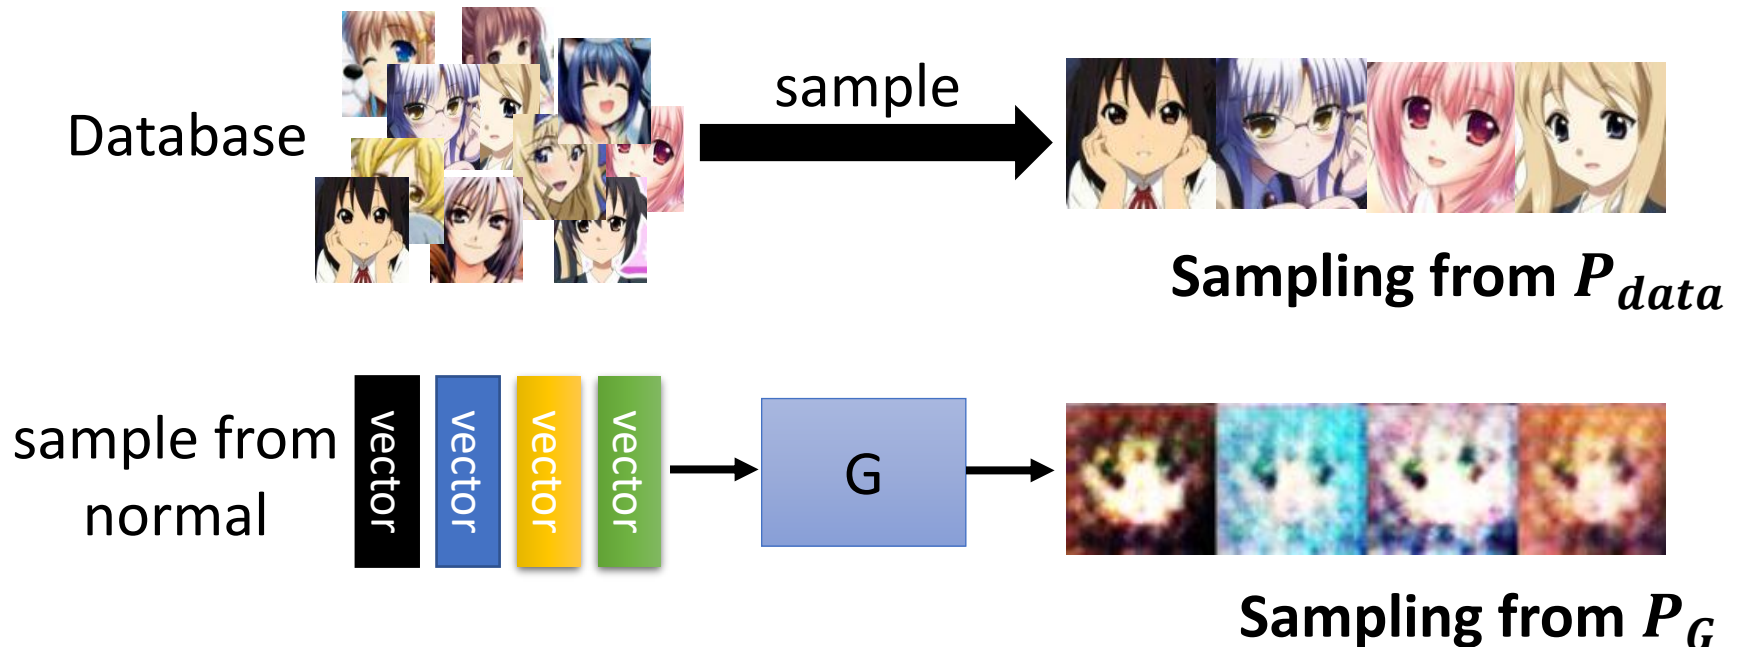
Divergence between distributions P_G and P_{data}

How to compute the divergence?

Sampling is good enough


$$G^* = \arg \min_G \text{Div}(P_G, P_{data})$$

Although we do not know the distributions of P_G and P_{data} , we can sample from them.



Discriminator $G^* = \arg \min_G \text{Div}(P_G, P_{data})$

★ : data sampled from P_{data} ★ : data sampled from P_G



Training: $D^* = \arg \max_D V(D, G)$ The value is related to JS divergence.

Objective Function for D

$$V(G, D) = E_{y \sim P_{data}} [\log D(y)] + E_{y \sim P_G} [\log(1 - D(y))]$$

$D^* = \arg \max_D V(D, G)$ = Training classifier:
negative cross entropy minimize cross entropy

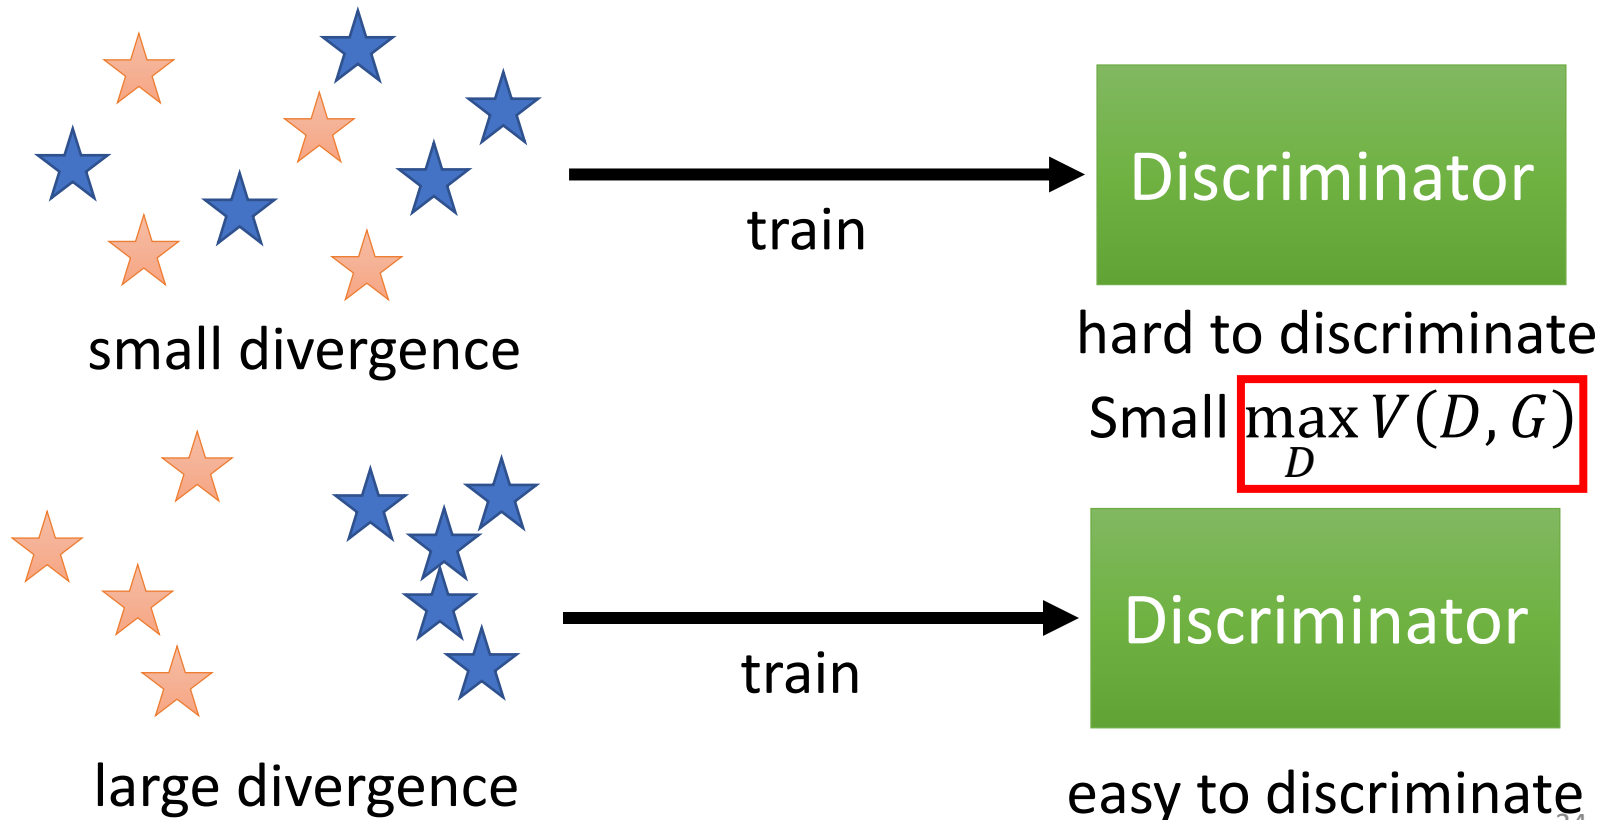
Discriminator $G^* = \arg \min_G \text{Div}(P_G, P_{data})$

★ : data sampled from P_{data}

★ : data sampled from P_G

Training:

$$D^* = \arg \max_D V(D, G)$$




$$G^* = \arg \min_G \max_D V(G, D)$$

$$D^* = \arg \max_D V(D, G)$$

The maximum objective value is related to JS divergence.

- Initialize generator and discriminator
- In each training iteration:

Step 1: Fix generator G , and update discriminator D

Step 2: Fix discriminator D , and update generator G

Can we use other divergence?

Name	$D_f(P\ Q)$	Generator $f(u)$
Total variation	$\frac{1}{2} \int p(x) - q(x) \, dx$	$\frac{1}{2} u - 1 $
Kullback-Leibler	$\int p(x) \log \frac{p(x)}{q(x)} \, dx$	$u \log u$
Reverse Kullback-Leibler	$\int q(x) \log \frac{q(x)}{p(x)} \, dx$	$-\log u$
Pearson χ^2	$\int \frac{(q(x)-p(x))^2}{p(x)} \, dx$	$(u - 1)^2$
Neyman χ^2	$\int \frac{(p(x)-q(x))^2}{q(x)} \, dx$	$\frac{(1-u)^2}{u}$
Squared Hellinger	$\int \left(\sqrt{p(x)} - \sqrt{q(x)} \right)^2 \, dx$	$(\sqrt{u} - 1)^2$
Jeffrey	$\int (p(x) - q(x)) \log \left(\frac{p(x)}{q(x)} \right) \, dx$	$(u - 1) \log u$
Jensen-Shannon	$\frac{1}{2} \int p(x) \log \frac{2p(x)}{p(x)+q(x)} + q(x) \log \frac{2q(x)}{p(x)+q(x)} \, dx$	$-(u + 1) \log \frac{1+u}{2} + u \log u$
Jensen-Shannon-weighted	$\int p(x) \pi \log \frac{p(x)}{\pi p(x) + (1-\pi)q(x)} + (1 - \pi)q(x) \log \frac{q(x)}{\pi p(x) + (1-\pi)q(x)} \, dx$	$\pi u \log u - (1 - \pi + \pi u) \log(1 - \pi + \pi u)$
GAN	$\int p(x) \log \frac{2p(x)}{p(x)+q(x)} + q(x) \log \frac{2q(x)}{p(x)+q(x)} \, dx - \log(4)$	$u \log u - (u + 1) \log(u + 1)$

Name	Conjugate $f^*(t)$
Total variation	t
Kullback-Leibler (KL)	$\exp(t - 1)$
Reverse KL	$-1 - \log(-t)$
Pearson χ^2	$\frac{1}{4}t^2 + t$
Neyman χ^2	$2 - 2\sqrt{1 - t}$
Squared Hellinger	$\frac{t}{1-t}$
Jeffrey	$W(e^{1-t}) + \frac{1}{W(e^{1-t})} + t - 2$
Jensen-Shannon	$-\log(2 - \exp(t))$
Jensen-Shannon-weighted	$(1 - \pi) \log \frac{1-\pi}{1-\pi e^{t/\pi}}$
GAN	$-\log(1 - \exp(t))$

Using the divergence
you like ☺

<https://arxiv.org/abs/1606.00709>

GAN is difficult to train

NO PAIN

NO GAN

(I found this joke from 陳柏文's facebook.)



Tips for GAN

JS divergence is not suitable

- In most cases, P_G and P_{data} are not overlapped.
- 1. The nature of data


Both P_{data} and P_G are low-dim manifold in high-dim space.

The overlap can be ignored.

- 2. Sampling

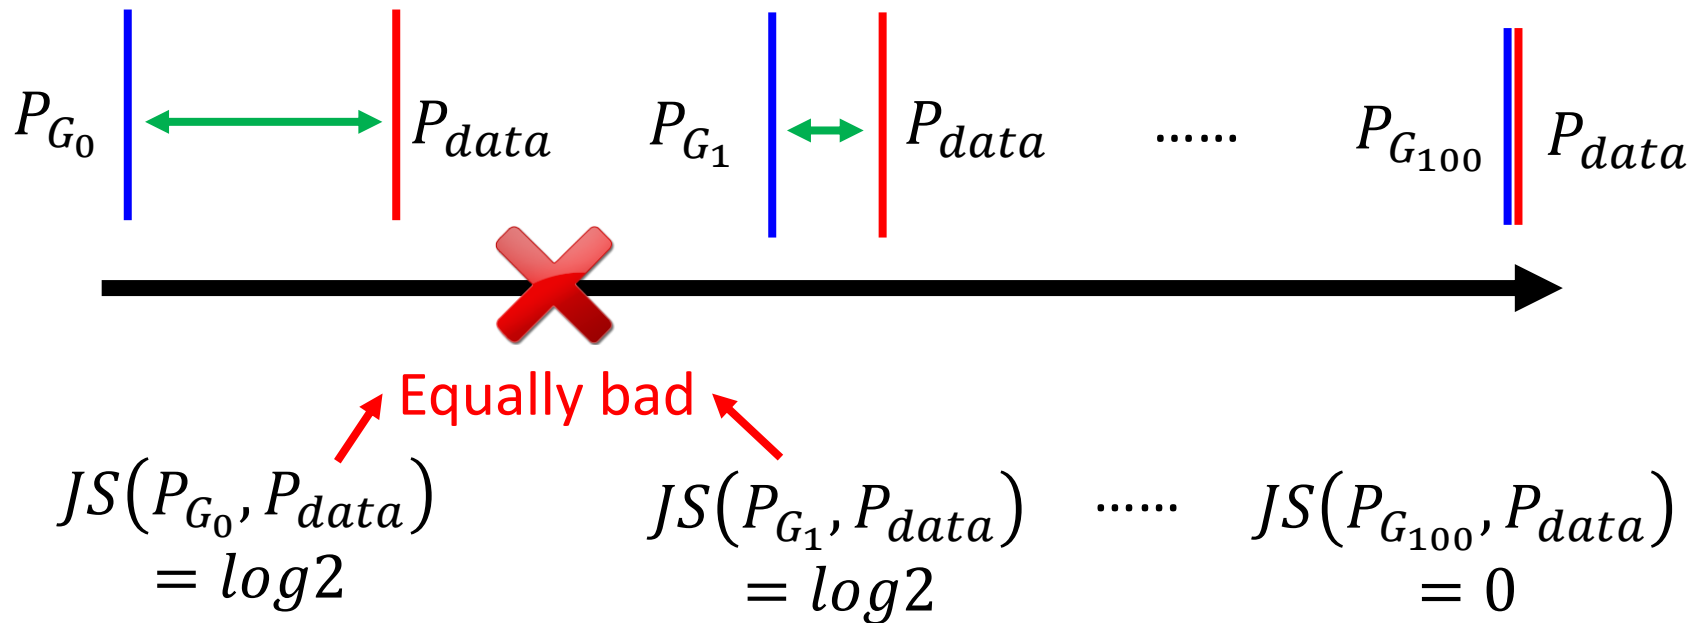
Even though P_{data} and P_G have overlap.

If you do not have enough sampling



What is the problem of JS divergence?

JS divergence is always $\log 2$ if two distributions do not overlap.




Intuition: If two distributions do not overlap, binary classifier achieves 100% accuracy.

Its accuracy (or loss) means nothing during GAN training.

Wasserstein distance


- Considering one distribution P as a pile of earth, and another distribution Q as the target
- The average distance the earth mover has to move the earth.



$$W(P, Q) = d$$




Wasserstein distance



There are many possible “moving plans”.

Using the “moving plan” with the smallest average distance to define the Wasserstein distance.

What is the problem of JS divergence?



$$JS(P_{G_0}, P_{data}) = \log 2$$

$$JS(P_{G_1}, P_{data}) = \log 2$$

$$JS(P_{G_{100}}, P_{data}) = 0$$


$$W(P_{G_0}, P_{data}) = d_0$$

$$W(P_{G_1}, P_{data}) = d_1$$


$$W(P_{G_{100}}, P_{data}) = 0$$

Better!


What is the problem of JS divergence?



pigment spot
(limpet, *Patella*)



Complex eye
(octopus)



$$\max_{D \in 1\text{-Lipschitz}} \{E_{x \sim P_{data}}[D(x)] - E_{x \sim P_G}[D(x)]\}$$

- Original WGAN → Weight


Force the parameters w between c and $-c$

After parameter update, if $w > c$, $w = c$; if $w < -c$, $w = -c$

- Improved WGAN → Gradient Penalty

Keep the gradient close to 1

<https://arxiv.org/abs/1704.00028>




- Spectral Normalization → Keep gradient norm smaller than 1 everywhere

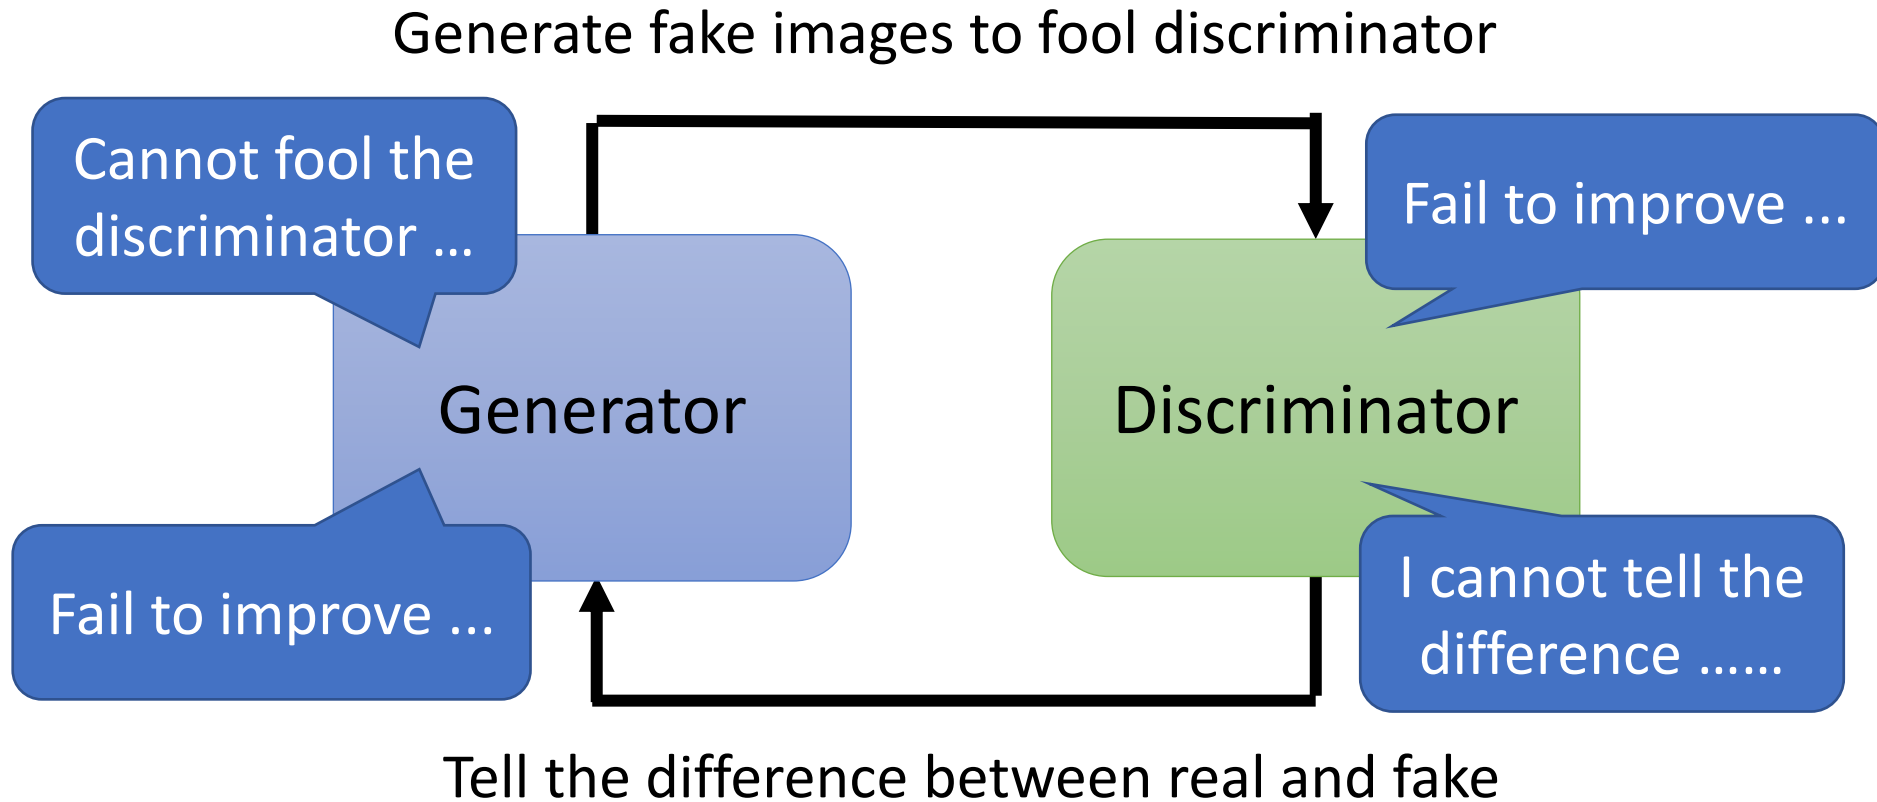
<https://arxiv.org/abs/1802.05957>

GAN is still challenging ...

- Generator and Discriminator needs to match each other (棋逢敵手)



GAN is still challenging ...



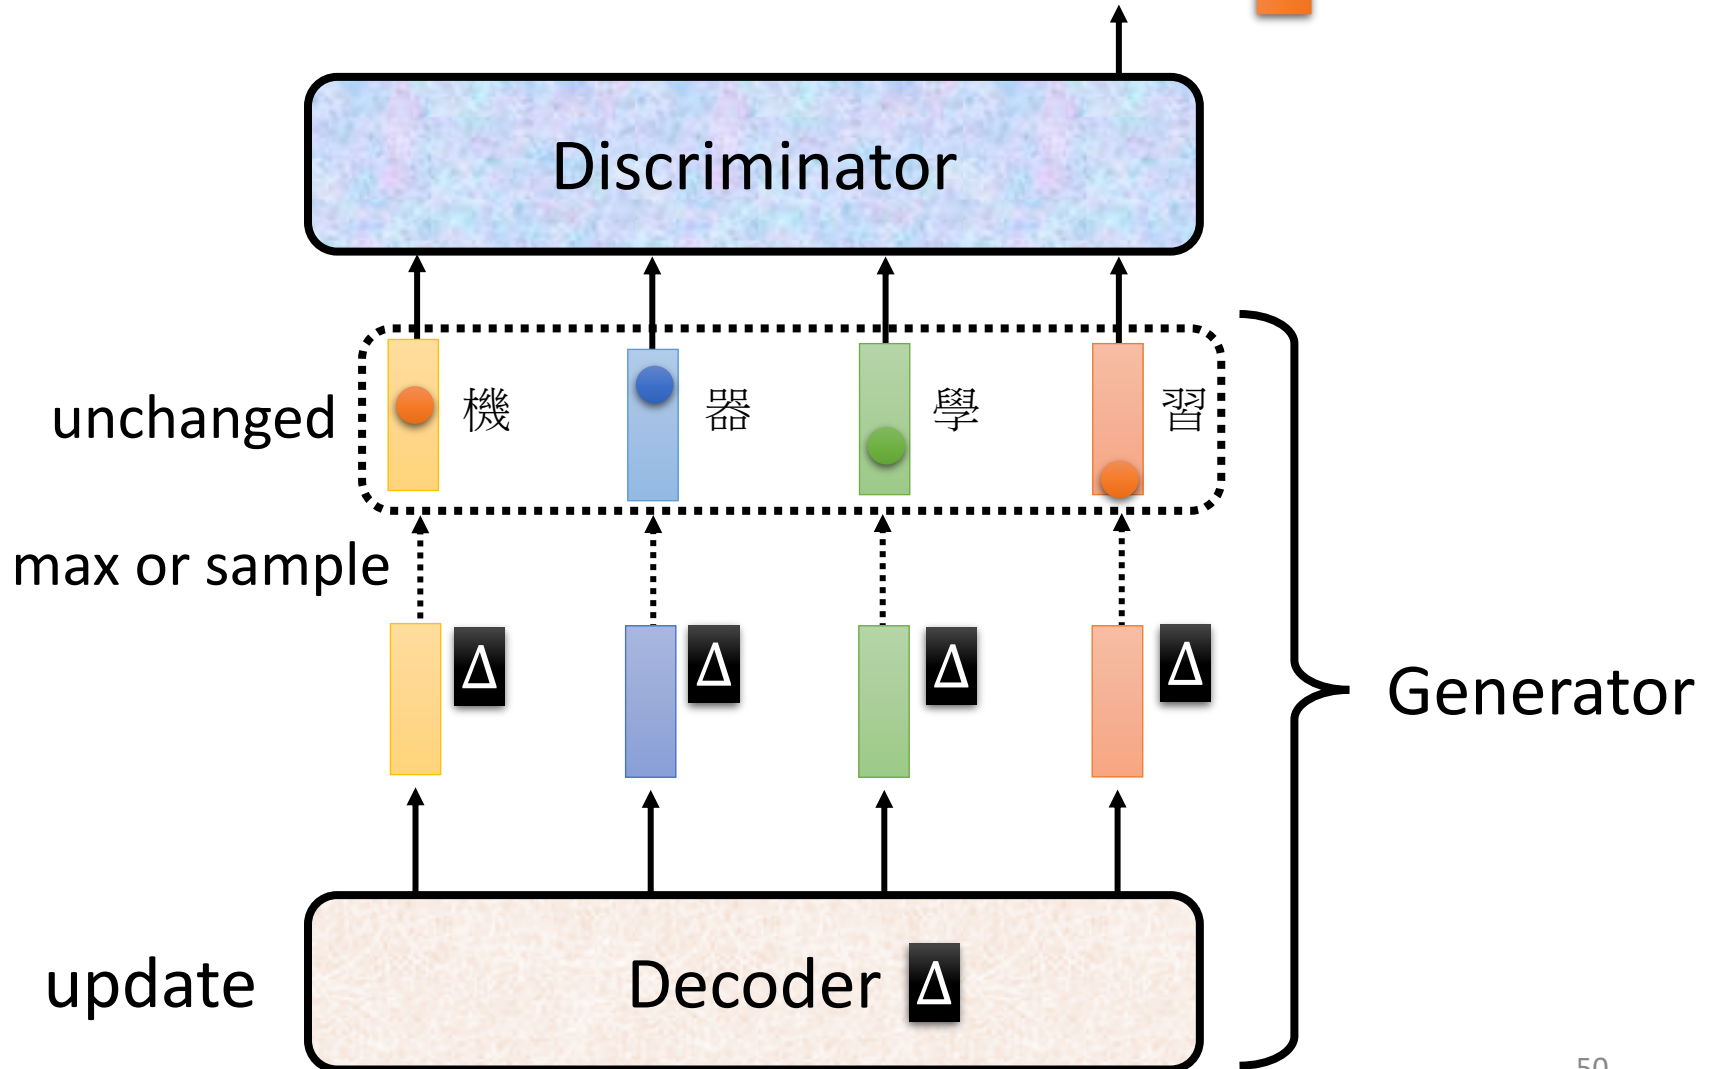
More Tips

- Tips from Soumith
 - <https://github.com/soumith/ganhacks>
- Tips in DCGAN: Guideline for network architecture design for image generation
 - <https://arxiv.org/abs/1511.06434>
- Improved techniques for training GANs
 - <https://arxiv.org/abs/1606.03498>
- Tips from BigGAN
 - <https://arxiv.org/abs/1809.11096>

GAN for Sequence Generation

Non-differentiable ...

score  unchanged




GAN for Sequence Generation

Reinforcement learning (RL) is involved

RL is difficult to train

GAN is difficult to train




Sequence Generation GAN (RL+GAN)

GAN for Sequence Generation

- Usually, the generator are fine-tuned from a model learned by other approaches.
- However, with enough hyperparameter-tuning and tips, ScarchGAN can train from scratch.

Training language
GANs from Scratch

[https://arxiv.org/abs/
1905.09922](https://arxiv.org/abs/1905.09922)



Generative Models

- This lecture: Generative Adversarial Network (GAN)




Full version

https://www.youtube.com/playlist?list=PLJV_el3uVTsMq6JFPW35BCiOQTsoqwNw

More Generative Models

Variational
Autoencoder (VAE)



<https://youtu.be/8zomhgKrmQ>

FLOW-based
Model



<https://youtu.be/uXY18nzdSsM>


Possible Solution?



$$\begin{bmatrix} 0.3 \\ -0.1 \\ \vdots \\ -0.7 \end{bmatrix}$$




$$\begin{bmatrix} 0.1 \\ -0.1 \\ \vdots \\ 0.7 \end{bmatrix}$$




$$\begin{bmatrix} -0.3 \\ 0.1 \\ \vdots \\ 0.9 \end{bmatrix}$$



$$\begin{bmatrix} 0.7 \\ 0.1 \\ \vdots \\ -0.9 \end{bmatrix}$$



$$\begin{bmatrix} -0.1 \\ 0.8 \\ \vdots \\ 0.8 \end{bmatrix}$$



Using typical learning approaches?

Generative Latent Optimization (GLO), <https://arxiv.org/abs/1707.05776>

Gradient Origin Networks, <https://arxiv.org/abs/2007.02798>

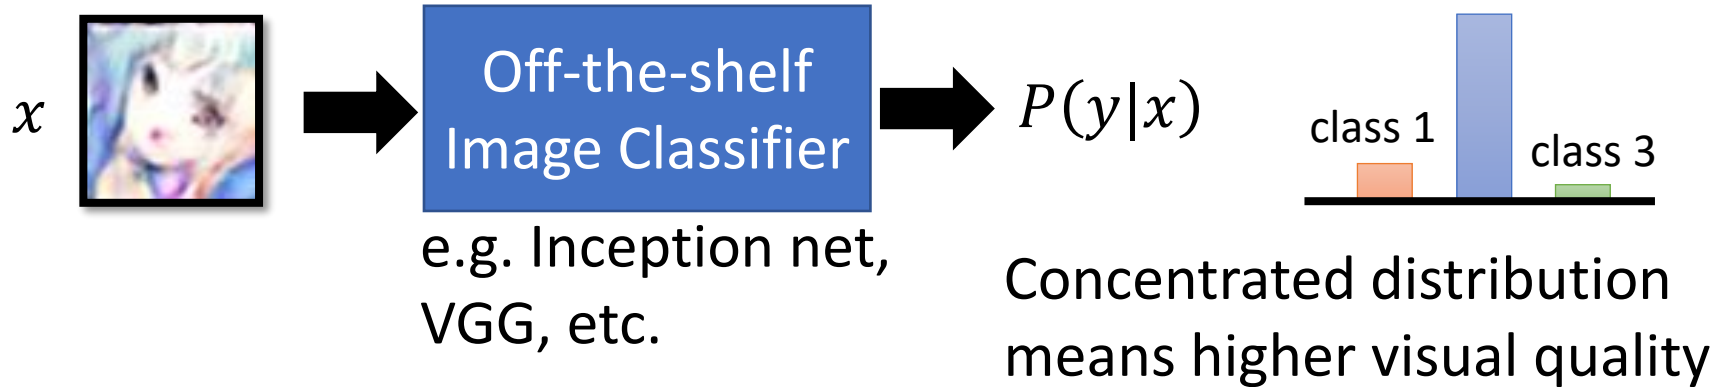


Evaluation of Generation

Inception Score

x : image


y : class (output of CNN)





Mode Collapse

★ : real data

★ : generated data



Mode Dropping



Generator switches mode during training

Generator
at iteration t



Generator
at iteration $t+1$




Generator
at iteration $t+2$





Mode
easy to








$$P(y) = \frac{1}{N} \sum_n P(y^n | x^n)$$



Uniform distribution
means higher variety

60

Inception Score



Inception Score

[Tim Salimans, et al., NIPS 2016]


$$= \sum_x \sum_y \underline{P(y|x) \log P(y|x)} - \underline{\sum_y P(y) \log P(y)}$$

Negative entropy of $P(y|x)$


Entropy of $P(y)$

Fréchet Inception Distance (FID)

<https://arxiv.org/pdf/1706.08500.pdf>



GAN	DISCRIMINATOR LOSS	GENERATOR LOSS
MM GAN	$\mathcal{L}_D^{\text{GAN}} = -\mathbb{E}_{x \sim p_d} [\log(D(x))] + \mathbb{E}_{\hat{x} \sim p_g} [\log(1 - D(\hat{x}))]$	$\mathcal{L}_G^{\text{GAN}} = -\mathcal{L}_D^{\text{GAN}}$
NS GAN	$\mathcal{L}_D^{\text{NSGAN}} = \mathcal{L}_D^{\text{GAN}}$	$\mathcal{L}_G^{\text{NSGAN}} = \mathbb{E}_{\hat{x} \sim p_g} [\log(D(\hat{x}))]$
WGAN	$\mathcal{L}_D^{\text{WGAN}} = -\mathbb{E}_{x \sim p_d} [D(x)] + \mathbb{E}_{\hat{x} \sim p_g} [D(\hat{x})]$	$\mathcal{L}_G^{\text{WGAN}} = -\mathcal{L}_D^{\text{WGAN}}$
WGAN GP	$\mathcal{L}_D^{\text{WGAN}} = \mathcal{L}_D^{\text{WGAN}} + \lambda \mathbb{E}_{\hat{x} \sim p_g} [(\nabla D(\alpha x + (1 - \alpha)\hat{x}) _2 - 1)^2]$	$\mathcal{L}_G^{\text{WGAN}} = -\mathbb{E}_{\hat{x} \sim p_g} [D(\hat{x})]$
LS GAN	$\mathcal{L}_D^{\text{LSGAN}} = -\mathbb{E}_{x \sim p_d} [(D(x) - 1)^2] + \mathbb{E}_{\hat{x} \sim p_g} [D(\hat{x})^2]$	$\mathcal{L}_G^{\text{LSGAN}} = -\mathbb{E}_{\hat{x} \sim p_g} [(D(\hat{x}) - 1)^2]$
DRAGAN	$\mathcal{L}_D^{\text{DRAGAN}} = \mathcal{L}_D^{\text{GAN}} + \lambda \mathbb{E}_{\hat{x} \sim p_d + \mathcal{N}(0, c)} [(\nabla D(\hat{x}) _2 - 1)^2]$	$\mathcal{L}_G^{\text{DRAGAN}} = -\mathcal{L}_D^{\text{NSGAN}}$
BEGAN	$\mathcal{L}_D^{\text{BEGAN}} = \mathbb{E}_{x \sim p_d} [x - \text{AE}(x) _1] - k_t \mathbb{E}_{\hat{x} \sim p_g} [\hat{x} - \text{AE}(\hat{x}) _1]$	$\mathcal{L}_G^{\text{BEGAN}} = \mathbb{E}_{\hat{x} \sim p_g} [\hat{x} - \text{AE}(\hat{x}) _1]$





FID: Smaller is better

Are GANs Created Equal? A Large-Scale Study

<https://arxiv.org/abs/1711.10337>

We don't want memory GAN.

- Using k-nearest neighbor to check whether the generator generates new objects



To learn more about evaluation ...

	Measure	Description
Quantitative	1. Average Log-likelihood [18, 22]	• Log likelihood of explaining realworld held out/test data using a density estimated from the generated data (e.g. using KDE or Parzen window estimation). $L = \frac{1}{N} \sum_i \log P_{model}(\mathbf{x}_i)$
	2. Coverage Metric [33]	• The probability mass of the true data "covered" by the model distribution $C := P_{data}(dP_{model} > t)$ with t such that $P_{model}(dP_{model} > t) = 0.95$
	3. Inception Score (IS) [3]	• KLD between conditional and marginal label distributions over generated data. $\exp(\mathbb{E}_{\mathbf{x}}[\mathbb{KL}(p(y \mathbf{x}) p(y))])$
	4. Modified Inception Score (m-IS) [34]	• Encourages diversity within images sampled from a particular category. $\exp(\mathbb{E}_{\mathbf{x}_i}[\mathbb{E}_{\mathbf{x}_j}[(\mathbb{KL}(P(y \mathbf{x}_i)) P(y \mathbf{x}_j))]])$
	5. Mode Score (MS) [35]	• Similar to IS but also takes into account the prior distribution of the labels over real data. $\exp(\mathbb{E}_{\mathbf{x}}[\mathbb{KL}(p(y \mathbf{x}) p(y^{train}))]) - \mathbb{KL}(p(y) p(y^{train}))$
	6. AM Score [36]	• Takes into account the KLD between distributions of training labels vs. predicted labels, as well as the entropy of predictions. $\mathbb{KL}(p(y^{train}) p(y)) + \mathbb{E}_{\mathbf{x}}[H(y \mathbf{x})]$
	7. Fréchet Inception Distance (FID) [37]	• Wasserstein-2 distance between multi-variate Gaussians fitted to data embedded into a feature space $FID(r, g) = \ \mu_r - \mu_g\ _2^2 + \text{Tr}(\Sigma_r + \Sigma_g - 2(\Sigma_r \Sigma_g)^{\frac{1}{2}})$
	8. Maximum Mean Discrepancy (MMD) [38]	• Measures the dissimilarity between two probability distributions P_r and P_g using samples drawn independently from each distribution. $M_k(P_r, P_g) = \mathbb{E}_{\mathbf{x}, \mathbf{x}' \sim P_r}[k(\mathbf{x}, \mathbf{x}')] - 2\mathbb{E}_{\mathbf{x} \sim P_r, \mathbf{y} \sim P_g}[k(\mathbf{x}, \mathbf{y})] + \mathbb{E}_{\mathbf{y}, \mathbf{y}' \sim P_g}[k(\mathbf{y}, \mathbf{y}')]$
	9. The Wasserstein Critic [39]	• The critic (e.g. an NN) is trained to produce high values at real samples and low values at generated samples $W(\mathbf{x}_{test}, \mathbf{x}_g) = \frac{1}{N} \sum_{i=1}^N \hat{f}(\mathbf{x}_{test}[i]) - \frac{1}{N} \sum_{i=1}^N \hat{f}(\mathbf{x}_g[i])$
	10. Birthday Paradox Test [27]	• Measures the support size of a discrete (continuous) distribution by counting the duplicates (near duplicates)
	11. Classifier Two Sample Test (C2ST) [40]	• Answers whether two samples are drawn from the same distribution (e.g. by training a binary classifier)
	12. Classification Performance [1, 15]	• An indirect technique for evaluating the quality of unsupervised representations (e.g. feature extraction; FCN score). See also the GAN Quality Index (GQI) [41].
	13. Boundary Distortion [42]	• Measures diversity of generated samples and covariate shift using classification methods.
	14. Number of Statistically-Different Bins (NDB) [43]	• Given two sets of samples from the same distribution, the number of samples that fall into a given bin should be the same up to sampling noise
	15. Image Retrieval Performance [44]	• Measures the distributions of distances to the nearest neighbors of some query images (i.e. diversity)
	16. Generative Adversarial Metric (GAM) [31]	• Compares two GANs by having them engaged in a battle against each other by swapping discriminators or generators. $p(\mathbf{x} y=1; M_1^*)/p(\mathbf{x} y=1; M_2^*) = (p(y=1 \mathbf{x}; D_1)p(\mathbf{x}; G_2))/(p(y=1 \mathbf{x}; D_2)p(\mathbf{x}; G_1))$
	17. Tournament Win Rate and Skill Rating [45]	• Implements a tournament in which a player is either a discriminator that attempts to distinguish between real and fake data or a generator that attempts to fool the discriminators into accepting fake data as real.
	18. Normalized Relative Discriminative Score (NRDS) [32]	• Compares n GANs based on the idea that if the generated samples are closer to real ones, more epochs would be needed to distinguish them from real samples.
	19. Adversarial Accuracy and Divergence [46]	• Adversarial Accuracy: Computes the classification accuracies achieved by the two classifiers, one trained on real data and another on generated data, on a labeled validation set to approximate $P_g(y \mathbf{x})$ and $P_r(y \mathbf{x})$. Adversarial Divergence: Computes $\mathbb{KL}(P_g(y \mathbf{x}) P_r(y \mathbf{x}))$
	20. Geometry Score [47]	• Compares geometrical properties of the underlying data manifold between real and generated data.
	21. Reconstruction Error [48]	• Measures the reconstruction error (e.g. L_2 norm) between a test image and its closest generated image by optimizing for z (i.e. $\min_z \ G(\mathbf{z}) - \mathbf{x}^{(test)}\ _2^2$)
	22. Image Quality Measures [49, 50, 51]	• Evaluates the quality of generated images using measures such as SSIM, PSNR, and sharpness difference
	23. Low-level Image Statistics [52, 53]	• Evaluates how similar low-level statistics of generated images are to those of natural scenes in terms of mean power spectrum, distribution of random filter responses, contrast distribution, etc.
	24. Precision, Recall and F_1 score [23]	• These measures are used to quantify the degree of overfitting in GANs, often over toy datasets.
Qualitative	1. Nearest Neighbors	• To detect overfitting, generated samples are shown next to their nearest neighbors in the training set
	2. Rapid Scene Categorization [18]	• In these experiments, participants are asked to distinguish generated samples from real images in a short presentation time (e.g. 100 ms); i.e. real v.s. fake
	3. Preference Judgment [54, 55, 56, 57]	• Participants are asked to rank models in terms of the fidelity of their generated images (e.g. pairs, triples)
	4. Mode Drop and Collapse [58, 59]	• Over datasets with known modes (e.g. a GMM or a labeled dataset), modes are computed as by measuring the distances of generated data to mode centers
	5. Network Internals [1, 60, 61, 62, 63, 64]	• Regards exploring and illustrating the internal representation and dynamics of models (e.g. space continuity) as well as visualizing learned features

Pros and cons of GAN evaluation measures

<https://arxiv.org/abs/1802.03446>

Concluding Remarks

Introduction of Generative Models

Generative Adversarial Network (GAN)

Theory behind GAN

Tips for GAN

Evaluation of Generative Models

